

# High-Accuracy Micro-Defect Detection System

A state-of-the-art computer vision pipeline for detecting micro-defects ("blisters") on mobile phone backglass with >99% recall and >95% precision.

## Overview

This project implements a **Three-Stage Refinement Pipeline** that progressively narrows focus from raw image to precise defect detection:

1. **Stage 1:** ROI Isolation - Find and align the phone backglass
2. **Stage 2:** Candidate Detection - High-recall search for all potential defects
3. **Stage 3:** Verification - High-precision filtering and final segmentation

## Architecture

### Stage 1: Robust ROI Isolation

- **Model:** YOLOv8-Small + Segment Anything Model (SAM)
- **Purpose:** Extract perfectly aligned backglass region
- **Output:** Standardized, rectangular backglass image

### Stage 2: High-Recall Candidate Search

- **Model:** YOLOv8-Large-Seg with P2 head (stride-4 detection)
- **Strategy:** Adaptive SAHI (Slicing Aided Hyper Inference)
- **Purpose:** Find every possible defect candidate
- **Output:** List of candidates with high recall, some false positives

### Stage 3: High-Precision Verification

- **Model:** Lightweight U-Net Refiner
- **Purpose:** Filter false positives, generate clean masks
- **Output:** Final defect list with bounding boxes and segmentation masks

# Project Phases

## Phase 1: Data Foundation (Weeks 1-4)

- **Environment Setup:** GPU server with PyTorch, OpenCV, CUDA
- **Image Acquisition:** 2,000+ high-resolution images
- **Annotation:** Pixel-perfect segmentation masks for all defects
- **Augmentation Library:** 300+ defect patches for copy-paste augmentation

## Phase 2: Model Development (Weeks 5-8)

- **Stage 1 Training:** YOLOv8-Small on ROI bounding boxes
- **Stage 2 Training:** P2-enabled YOLOv8-Large-Seg with heavy augmentation
- **Stage 3 Training:** U-Net on balanced true/false positive crops

## Phase 3: Integration & Validation (Weeks 9-10)

- **Pipeline Integration:** End-to-end inference script
- **Performance Benchmarking:** Validation on held-out test set
- **Error Analysis:** Systematic failure case analysis

## Phase 4: Finalization (Weeks 11-12)

- **Iterative Improvement:** Retrain based on error analysis
- **Documentation:** Complete technical documentation
- **Demo Application:** Interactive visualization tool

# Technical Specifications

## Hardware Requirements

- **GPU:** NVIDIA A6000/RTX 4090 (24GB+ VRAM recommended)
- **RAM:** 32GB+ system memory
- **Storage:** 500GB+ for datasets and models

## Software Dependencies

```
# Core ML Framework  
torch>=2.0.0  
torchvision>=0.15.0  
ultralytics>=8.0.0
```

```
# Computer Vision  
opencv-python>=4.8.0  
segment-anything>=1.0
```

```
# Data Processing  
numpy>=1.24.0  
pandas>=2.0.0  
pillow>=10.0.0
```

```
# Visualization  
matplotlib>=3.7.0  
seaborn>=0.12.0
```

## Model Configurations

### Stage 1: ROI Detector

```
# Training Parameters  
imgsz: 640  
batch: 16  
epochs: 50  
lr0: 0.001  
optimizer: AdamW  
patience: 10
```

### Stage 2: Candidate Detector

```
# Training Parameters  
imgsz: 1280  
batch: 4  
epochs: 150  
lr0: 0.001  
weight_decay: 0.0005
```

```
# Augmentation  
mosaic: 1.0  
fliplr: 0.5  
flipud: 0.5  
copy_paste: 0.5
```

```
# Loss Weights  
box: 7.5  
cls: 0.5  
dfl: 1.5  
fl_gamma: 1.5
```

## Stage 3: Refiner U-Net

*# Architecture*

**input\_size:** 128

**channels:** [16, 32, 64]

**depth:** 3

*# Training*

**batch\_size:** 64

**epochs:** 75

**lr:** 3e-4

**optimizer:** Adam

*# Loss Function*

**loss:** 0.7 \* DiceLoss + 0.3 \* BCELoss

## Performance Targets

| Metric           | Target | Description                                    |
|------------------|--------|------------------------------------------------|
| Defect Recall    | >99.0% | Percentage of true defects detected            |
| Defect Precision | >95.0% | Percentage of detections that are real defects |
| Mask IoU         | >0.70  | Segmentation mask accuracy                     |
| ROI Success Rate | >99.9% | Backglass isolation success rate               |

## Usage

### Training Pipeline

*# Stage 1: Train ROI detector*

python train\_roi.py --data roi\_dataset.yaml --epochs 50

*# Stage 2: Train candidate detector*

python train\_candidates.py --data defect\_dataset.yaml --epochs 150

*# Stage 3: Train refiner*

python train\_refiner.py --data refiner\_dataset.yaml --epochs 75

## Inference Pipeline

*# Run complete detection pipeline*

```
python detect.py --image path/to/image.jpg --output results.json
```

*# Batch processing*

```
python detect.py --folder path/to/images/ --output results/
```

## Output Format

```
{
  "image": "sample.jpg",
  "defects": [
    {
      "id": 1,
      "bbox": [x, y, width, height],
      "mask": "base64_encoded_mask",
      "confidence": 0.95,
      "area_pixels": 156
    }
  ],
  "total_defects": 1,
  "processing_time": 1.2
}
```

## Key Innovations

### P2 Feature Head

- Enables detection at stride-4 for objects as small as 8-16 pixels
- Critical architectural modification for micro-defect detection

### Adaptive SAHI Strategy

- Dynamic tile sizing based on defect probability
- Focuses computational resources on high-likelihood regions

### Specialized Refiner Network

- Trained exclusively on hard negatives from Stage 2
- Acts as expert classifier for final verification

## Copy-Paste Augmentation

- Library of 300+ defect patches with transparent backgrounds
- Realistic blending with Poisson editing techniques

## Data Requirements

### Image Collection

- **Quantity:** 2,000+ high-resolution images
- **Variety:** All defect types, sizes, locations
- **Quality:** Consistent lighting, sharp focus
- **Negatives:** Clean units + confusing artifacts (dust, fibers)

### Annotation Standards

- **Tool:** CVAT or Labelbox recommended
- **Format:** COCO or YOLO segmentation format
- **Quality:** Pixel-perfect masks, consistent labeling
- **Validation:** Multi-annotator agreement >95%

## Deployment Considerations

### Inference Optimization

- **Model Quantization:** INT8 quantization for production
- **TensorRT:** GPU acceleration for real-time processing
- **Batch Processing:** Optimize for throughput vs latency

### Quality Assurance

- **Active Learning:** Continuous model improvement loop
- **Hard Negative Mining:** Systematic false positive collection
- **Performance Monitoring:** Real-time accuracy tracking

## Troubleshooting

### Common Issues

1. **Low Recall:** Check P2 head implementation, increase augmentation

2. **High False Positives:** Retrain Stage 3 with more hard negatives
3. **ROI Failures:** Improve Stage 1 dataset diversity
4. **Memory Issues:** Reduce batch size, use gradient accumulation

## Performance Optimization

- Use mixed precision training (FP16)
- Implement gradient checkpointing for large models
- Optimize data loading with multiple workers

## Contributing

### Development Workflow

1. Fork repository and create feature branch
2. Implement changes with comprehensive tests
3. Update documentation and performance benchmarks
4. Submit pull request with detailed description

### Code Standards

- Follow PEP 8 style guidelines
- Include docstrings for all functions
- Add unit tests for new functionality
- Maintain >90% code coverage

## License

This project is proprietary and confidential. All rights reserved.

## Contact

For technical questions or implementation support, please contact the development team.

---

**Note:** This system represents cutting-edge computer vision technology specifically designed for industrial quality control applications. The three-stage architecture ensures maximum accuracy while maintaining practical deployment feasibility.